

# Constraints, governance and risk

Entirely fictional. Case snapshot: August 31, 2026. Synthetic organizational evidence for a hackathon prototype.

## Policy status

These are fictional Asteron pilot rules, not a statement of current law. Privacy, cross-border processing, customer contracts, employment implications and any sector-specific obligations must be reviewed by the relevant legal and security owners before real deployment. The hackathon uses synthetic data only.

## Customer data and security

Collect only necessary fields. Remove secrets, access tokens, sensitive attachments and unnecessary personal data before ingestion. Tenant-level access controls must apply during retrieval and generation. A user must never see another customer's material. Use an approved isolated environment with contractual restrictions on training from customer content. Keep encrypted storage and transfer, role-based access, and an audit trail. Pilot caches expire after 30 days; approved audit records after 90 days, subject to a documented legal hold.

## Technology and integration limits

CaseHarbor exports arrive nightly. CRM entitlements can therefore be a day old. The acquired product space lacks a production connector. The prototype must simulate authorization rather than claim live permission enforcement. Production connector work needs separate acceptance tests. No production action, customer email, billing adjustment, entitlement change or data replay is allowed from the hackathon.

## Operational limits

Pilot incremental spend ceiling: EUR 60,000 over eight weeks. Preserve existing customer commitments. Allow at most two hours per participating agent per week for training and feedback. Use reversible changes and a documented rollback. Measure extra review and correction time, not only gross time saved.

# Constraints, governance and risk / 2

## Human at the Center

Design around customers who need accurate, understandable help and employees who need usable support and room to exercise judgment. Include multilingual users and less experienced agents in testing. Offer a route to a person and a way to challenge an answer. Avoid turning quality monitoring into covert employee surveillance.

## Human in Command

The head of Support owns service outcomes, pilot scope, budget, capacity reassignment and stop decisions. Product owners authorize technical reference changes. Security controls access and incident escalation. Customer administrators authorize changes in their own environments. Accountability cannot be delegated to a model.

## Human in the Loop

A trained agent reviews customer-facing technical advice, checks applicability and confirms high-impact actions with the authorized owner. Human approval is required for data replay, security or permission changes, commercial commitments and uncertain cross-product remedies. Review must include enough source context and time to be meaningful.

## Human on the Loop

After separate validation, reversible internal tags or background issue clustering may run without per-item approval within explicit boundaries. A named operations owner reviews exceptions daily and samples at least 10% of outputs weekly. Provide audit logs and immediate disable/rollback. Stop on any cross-tenant exposure or unauthorized action; pause a capability if material error exceeds 5% in the weekly sample.

## Choosing the work design

Human-led work remains appropriate for empathy, relationship repair, business tradeoffs and accountability. AI-led / autonomous work requires bounded tasks, low consequence of error and effective intervention. Human-AI collaboration combines machine assistance with active human judgment. Non-AI organizational change may be sufficient where a policy, ownership, staffing or process decision is the real constraint. These are design choices, not a maturity ladder.