

Dataset guide | team only

Entirely fictional. Case snapshot: August 31, 2026. Evaluator material. Do not ingest into the prototype.

Start here

This folder contains 12 deliverables: eight PDFs and four CSVs. Everything about Asteron, its products, systems, employees, customers, costs and outcomes is synthetic. Dates span September 2025 through August 2026. This guide is team-only operational documentation and must stay outside the prototype's evidence index.

Safe ingestion

Upload exactly the nine files inside 01_Prototype_Inputs. Keep all three PDFs inside 02_Team_Only hidden. The full package ZIP is for the team; do not upload it to the prototype. The test-case PDF contains expected answers and is just as sensitive to evaluation validity as the ground-truth PDF.

Foundation sources

The Leverage Point PDF was available and used. The separate Human, AI & Human-AI Collaboration Guide was not available. Human roles follow the brief and the Leverage Point PDF. Before aligning to the missing guide verbatim, obtain it and review terminology. The original framework includes a support-knowledge example that could prime this case; use the example-free instruction on the next page for evaluation.

Suggested two-day sequence

Day 1: agree the audience and strategic question, ingest the allowlist, expose citations and build the reasoning flow. Day 2: test alternative rankings, intervention options, value and human roles. Run the five scenarios before showing the team answer key. Save the outputs and scores. For a fair comparison, give a simple use-case prompt and the systemic prompt the same evidence and model settings in separate fresh sessions.

Dataset guide | team only / 2

Example-free prototype instruction

Start from the company's strategic priorities. Follow Strategy -> Signals -> Patterns -> Root conditions -> Systemic connections -> Leverage points -> Value potential -> Right intervention -> Human-AI design -> Experiment. An observed complaint or bad KPI is a signal, not automatically a leverage point. Cite file names and record IDs. Separate observations, inference, assumptions and uncertainty.

Rank candidates before choosing technology

Score each candidate 1-5 on strategic relevance, systemic reach, causal depth, propagation potential, value multiplier, interventionability and evidence strength. Use equal weights by default and explain each score; report sensitivity to reasonable weighting changes. Do not include AI suitability in the leverage-point score. Ask what else would change if the candidate condition improved, and what could improve without changing it.

Choose and test the intervention

Compare Human-led, AI-led / autonomous, Human-AI collaboration and Non-AI organizational change. Evaluate AI suitability only after identifying the leverage point. Make Human at the Center, Human in Command, Human in the Loop and Human on the Loop explicit where relevant. Apply responsible-AI constraints throughout.

Value and output structure

Assess Efficiency & Capacity; Quality & Risk; Decision & Judgment; Innovation & Learning; Customer & Stakeholder Value; Capability & Agency; and Sustainable Growth. Include hard KPIs and ROI where defensible. Ask: If capacity is released, what happens to that capacity? Return a strategic objective, evidence-linked signals, competing hypotheses, candidate scores, counterfactuals, intervention choices, human roles, value assumptions and a falsifiable experiment. Never present correlation as proof of retention uplift.

Dataset guide | team only / 3

Prototype input inventory

1. Asteron_Company_Strategy_and_Context.pdf: business model, operating model, strategy and constraints. Future source: strategy documents, CRM and finance systems.
2. Asteron_Customer_Service_KPIs.csv: 48 rows, 12 months x three complexity segments plus all. Future source: ticketing analytics, workforce and finance data; account fields from CRM.
3. Asteron_Support_Tickets.csv: 50 closed, purposively selected cases. Future source: ticketing, customer interaction logs and engineering links.
4. Asteron_Customer_Feedback.csv: 20 comments linked to sampled cases. Future source: surveys, account reviews and CRM.
5. Asteron_Agent_Feedback.csv: 20 internal comments with role, date and optional ticket link. Future source: interviews, employee surveys and collaboration platforms.
6. Asteron_Customer_Service_Workflow.pdf: stages, branches, tools, teams and friction. Future source: process documentation and validated process-mining event logs.
7. Asteron_Knowledge_Base_Sample.pdf: 13 articles and two specialist notes with source and scope metadata. Future source: knowledge management and product operations spaces.
8. Asteron_AI_Opportunity_Backlog.pdf: nine sponsor proposals and unverified estimates. Future source: idea portfolio, business cases and pilot records.
9. Asteron_Constraints_Governance_and_Risk.pdf: fictional privacy, security, autonomy, operational and accountability rules. Future source: policies, contracts, architecture and risk registers.

Dataset guide | team only / 4

Team-only inventory

10. Asteron_Ground_Truth_TEAM_ONLY.pdf: causal answer key, evidence map, alternatives, illustrative scores, value economics and experiments. Hidden.

11. Asteron_Prototype_Test_Cases.pdf: five test questions, input subsets, expected answers and scoring rubrics. Hidden; copy only the question and authorized context per test.

12. Asteron_Dataset_Guide.pdf: this setup guide, ingestion rules and data dictionary. Hidden; only its example-free instruction may be copied.

How the files relate

month + complexity uniquely identifies KPI rows. ticket_id joins tickets to customer comments and optional agent comments. customer_id is a stable fictional account key; the sample has 23 accounts and does not represent all 650 customers. article_ids links tickets to KB-01 through KB-12; notes may reference SN-01. Ticket product values joined with + denote multi-product cases, not an extra product.

Sampling and join limitations

The 50 cases are a purposive teaching sample, not a random draw and not the population underlying the monthly KPI totals. Complex cases are deliberately represented more often. Do not recompute company rates from this sample or extrapolate its cause frequency. Customer comments are linked follow-ups, not independent corroborating observations. Dates in feedback identify the focal case date; employee-comment dates identify the comment date. No customer-level churn labels, revenue weights, event timestamps or active labor logs are supplied.

Process mining boundary

The workflow describes event order, but the case CSV contains aggregate contacts, handoffs and elapsed time. It is insufficient for real process mining. A future event log needs ticket ID, event timestamp, activity, accountable team, state transitions and waiting reasons. Do not fabricate those details from the workflow.

Dataset guide | team only / 5

KPI dictionary | grain and counts

All rows: synthetic=true; month is YYYY-MM; complexity is simple, medium, complex or all. Simple means one well-defined request with a standard path. Medium means configuration or coordination within a mainly single-product context. Complex means cross-product/integration investigation, material uncertainty or technical risk. Complexity is assigned after closure and is not the same as urgency.

Flow and performance denominators

cases_received counts newly opened requests in the month. cases_closed counts final closures in the month; backlog_open + cases_received - cases_closed = backlog_close. Reopens before final closure do not count as new arrivals. Segment classification is stable for this constructed roll-forward. fcr_cases, repeat_contact_cases and escalated_cases use cases_closed as denominator. FCR means one interaction cycle and no subsequent contact within seven days. Repeat contact means renewed customer contact within seven days after apparent resolution. A non-FCR case need not be a repeat case.

Surveys and rates

csat_responses counts returned 1-5 surveys; csat_positive_responses counts ratings 4 or 5. csat_pct = positives / responses x 100. fcr_pct, repeat_contact_pct and escalation_pct are their respective count / cases_closed x 100. All percentages are stored on a 0-100 scale. Survey response rates differ by complexity. Do not average segment percentages; aggregate their numerators and denominators. Seven-day outcomes for August closures are frozen on September 8, 2026; the business snapshot remains August 31.

Time, effort and cost

total_resolution_hours sums elapsed hours for closed cases; avg_resolution_hours divides by cases_closed. customer_wait_hours is the documented customer-wait portion of that elapsed total, not a labor measure. total_handoffs counts primary team changes; handoffs_per_case divides by cases_closed. support_cost_eur is an allocated fully loaded support cost pool attributed to closed cases, including support overhead but excluding separately budgeted engineering fixes. cost_per_closed_case_eur divides by cases_closed. It is not an incremental saving rate.

Dataset guide | team only / 6

KPI dictionary | capacity and accounts

`allocated_delivery_fte` is a 50/50/60 segment allocation summing to 160.

`closed_cases_per_delivery_fte` = `cases_closed` / `allocated_delivery_fte`. Allocation is held constant for comparability; local absences and shift overlap can still affect queues.

`defect_linked_cases` counts closed cases with an engineering defect link, without proving the defect caused the full delay.

CRM fields

`customers_open`, `customers_churned`, `customers_added` and `customers_close` appear only on all rows. Closing accounts = opening - churned + added; next month opens at prior closing.

`monthly_logo_churn_pct` = `churned` / `opening` x 100. `opening_cohort_remaining` tracks the surviving September 1, 2025 cohort. `annual_logo_retention_pct` appears only in August 2026: 528 / 600 = 88%. Blank earlier values mean insufficient 12-month cohort data, not zero. Do not add monthly churn rates to obtain annual retention.

Ticket dictionary

`ticket_id` and `customer_id` are identifiers; date is opening date YYYY-MM-DD. `region` and `customer_segment` describe the account. `contacts` includes the first interaction cycle; `handoffs` counts changes of primary team. `escalated` and `knowledge_article_used` use yes/no. `resolution_hours` includes waits and weekends. `csat_1_to_5` is an integer or blank for no response. `customer_sentiment` is a coarse synthetic annotation, not a classifier result. `resolution_outcome` is resolved, patched or workaround. A workaround means customer-accepted closure with continuing defect follow-up. `article_ids` is blank when no formal article was logged.

Feedback dictionary and missingness

Customer records contain `feedback_id`, `ticket_id`, `customer_id`, date, channel, sentiment, `csat_1_to_5` and comment. Agent records contain `feedback_id`, `employee_id`, role, region, date, optional `ticket_id` and comment. Blank ticket links mean general feedback. Employees are role-coded fictional respondents; no inference about named real people is intended. All four CSVs use UTF-8, comma delimiters, quoted text where necessary, one header row and blank missing values.

Verification and interpretation

Data checks reconcile backlog, account movement, segment totals, derived rates and identifiers. PDF pages were rendered and visually inspected. The adversarial review is a design assessment, not evidence that a model has passed a blind test. Use the five evaluation scenarios to measure that separately.