

Ground truth | team only

Entirely fictional. Case snapshot: August 31, 2026. Evaluator material. Do not ingest into the prototype.

Intended leverage point

Fragmented, inconsistent, and difficult-to-access organizational knowledge during complex customer interactions. More precisely, the intervention target is the relationship between reliable, version-specific organizational knowledge and the decisions made during cross-product service work. It is not simply a repository or an AI search feature.

Status of this answer key

This is the intended causal structure of a designed case, not a causal result proven by observational evidence. A strong engine should identify this as the leading testable hypothesis while considering release defects, routing metadata, customer waits, staffing overlap and incentives. It should not claim knowledge explains every case or that fixing it will deliver 93% retention.

Framework fidelity and limitations

The supplied Leverage Point PDF was read and used. The separate Human, AI & Human-AI Collaboration Guide could not be located. The brief and the Leverage Point PDF supply the human-role definitions used here. The source PDF has six scoring criteria in its final checklist; this brief explicitly adds propagation potential, so this package uses seven. AI suitability is excluded from that score.

Prevent answer leakage

Give the prototype only the nine files in 01_Prototype_Inputs. Keep this PDF, the test cases and the dataset guide outside its retrieval index, prompts and conversation. The original Leverage Point PDF includes essentially this support-knowledge example. Use only the example-free reasoning instruction in the dataset guide for an unprimed evaluation. Do not upload the full ZIP into the prototype.

Ground truth | team only / 2

Evidence map: observations that connect

KPIs: compare complex and simple rows across all 12 months. Complex work accumulates backlog, has more handoffs and lower first-contact resolution. May is an additional defect shock, not the whole pattern. These are symptoms and scale indicators, not root-cause labels.

Tickets AST-0004, 0021 and 0041: retry behavior and version checks recur across time. KB-03 supplies a general 60-second instruction; KB-12 supplies 120 seconds for managed adapters. Applicability, access and version context matter.

Tickets AST-0015, 0033 and 0046: replay decisions need a specialist check. KB-04 states a standard 24-hour window; SN-01 supplies an 8-hour exception for replay-v2. An escalation can be justified even when reference access improves.

Tickets AST-0007, 0025, 0036 and 0049: cross-product completion is interpreted differently. KB-08 describes final approval; KB-12 describes destination acknowledgment. Customer comments CF-03 and CF-17 show how the disagreement is experienced.

Tickets AST-0012, 0030 and 0043, with KB-10 and SN-02: inherited rule behavior and experienced-agent context recur. AST-0043 is a fast, successful complex case, useful counterevidence to the idea that complexity alone makes delay inevitable.

AF-06, 07, 08, 15, 18 and 20: review timing, experienced colleagues, separate spaces and applicability uncertainty appear in different roles. AF-10 suggests a reinforcing incentive condition: queue pressure crowds out maintenance.

Ground truth | team only / 3

Reasoning chain

Strategic priority -> retention, service consistency, profitable sustainable growth and productive use of released capacity.

Organizational signals -> repeated contacts, uncertain advice, specialist waits and cross-product interpretation differences.

Recurring patterns -> similar decision friction across retry, replay, event timing and inherited rules; some cases resolve without a code change once context is clarified.

Underlying condition -> reference material and tacit expertise do not consistently provide an accessible, authoritative answer for the specific tenant and version.

Systemic connections -> decision quality affects repeat work, escalation load, expert availability, onboarding, customer trust and capacity. Queue pressure can reduce maintenance, further increasing uncertainty.

Candidate leverage point -> improve how trusted, applicable knowledge reaches service decisions and how new learning returns to shared practice.

Value potential -> faster resolution and less rework, safer decisions, stronger agency, reusable learning, better customer experience and a plausible contribution to retention.

Intervention logic -> first clarify ownership, version/scope metadata, review triggers and escalation rules; then test AI assistance against a non-AI baseline.

Human-AI design -> AI retrieves, compares and flags conflicting sources; humans assess applicability, approve risky actions and remain accountable.

Experiment -> test matched complex scenarios across usual practice, governed references alone, and governed references plus AI.

Ground truth | team only / 4

Candidate comparison

Illustrative evaluator judgments, not measured scores. Scale 1-5, equal weights; maximum 35. Criteria order: strategic relevance, systemic reach, causal depth, propagation potential, value multiplier, interventionability, evidence strength.

Knowledge-to-decision reliability | 32/35

5, 5, 5, 5, 5, 4, 3. Strong cross-source coherence but a small curated sample and missing account-level linkage limit causal certainty.

Release quality and defect prevention | 27/35

5, 4, 4, 4, 4, 3, 3. Genuine upstream alternative. May incidents and AST-0038 support it; no-code resolutions across several other categories limit its explanatory reach.

Ownership, routing and staffing overlap | 25/35

4, 4, 3, 3, 4, 4, 3. Real delays from entitlement sync, language assignment and shift overlap; better assignment alone does not settle tenant-specific technical decisions.

Wrap-up writing effort | 18/35

3, 3, 1, 2, 2, 5, 2. A legitimate efficiency opportunity with measurable pilot evidence, but much weaker causal depth and propagation. The self-reported pilot is not a time-and-motion study.

Score discipline

Require an evidence citation and uncertainty note for each criterion. Strategic relevance concerns goals; systemic reach concerns connected parts; propagation concerns how a change spreads; value multiplier concerns distinct value dimensions. Avoid counting the same benefit four times. Test sensitivity to weights and missing evidence. An alternative ranking is acceptable if well supported; agreement with the key alone earns no credit.

Ground truth | team only / 5

Symptoms, root conditions and alternatives

Long resolution, low CSAT, repeated contacts and escalations are symptoms. Unclear source applicability, ownership boundaries and dependence on tacit expertise are candidate root conditions. Throughput incentives may reinforce them. Code defects, customer security approvals, stale contract data and coverage gaps are separate causal paths that need their own interventions.

Counterfactuals

If summaries became instant, agents would still need to decide which setting applies; little direct evidence connects wrap-up effort to contradictory advice. If knowledge access improved, a firewall approval or a query regression would still need the customer or Engineering. If routing alone improved, the correct team might still have to reconcile inherited behavior. Test whether version-correct advice improves before waiting times fall.

Tempting opportunities

Summaries, classification and drafting offer visible minutes saved. Routing and FAQ automation also have genuine merit. A weak use-case generator may rank the 750-hour summary extrapolation first, label all escalations waste, or propose a chatbot from CSAT alone. A strong engine separates opportunity attractiveness from systemic leverage and validates eligibility, overlap, review time and causal reach.

Adversarial review

One file can suggest a knowledge hypothesis: any authentic sample with exceptions can do that. No single file establishes its scale, strategic importance and causal reach. Agent comments include operational frustrations, and the KB does not label defects or a root problem. Release quality is a reasonable rival. The intended hypothesis is favored across all evidence because multiple recurring non-code issues connect reference applicability with decisions, handoffs and customer experience. This is an editorial review, not a blinded model benchmark; run the held-out tests before claiming discovery performance.

Ground truth | team only / 6

Intervention options

Human-led: paired diagnosis, coaching and explicit customer communication. AI-led / autonomous: only low-risk internal tags or pattern grouping under monitored boundaries. Human-AI collaboration: permission-aware retrieval, source comparison, synthesis and applicability checks with agent review. Non-AI organizational change: source owners, version tags, review triggers, protected learning time and clear cross-product ownership. Compare all four before choosing a combination.

Human roles

Human at the Center: customers and agents shape usefulness, accessibility and success measures. Human in Command: the head of Support owns the pilot and capacity redeployment; product owners approve technical knowledge. Human in the Loop: agents verify tenant applicability and customer advice; specialists approve high-impact actions. Human on the Loop: an operations owner monitors bounded background clustering, errors and exceptions. Monitoring is not a substitute for approval on risky actions.

Value across seven dimensions

Efficiency & Capacity: measure net active minutes saved after verification, maintenance and review; do not equate elapsed resolution hours with labor savings.

Quality & Risk: measure incorrect advice, unsafe actions, privacy failures and reopens. Decision & Judgment: assess whether agents select the correct version and explain uncertainty.

Innovation & Learning: count validated lessons incorporated and reused. Customer & Stakeholder Value: assess repeat-contact rate, CSAT and proactive-review usefulness.

Capability & Agency: assess agent confidence together with demonstrated accuracy and ability to override. Sustainable Growth: track consistency across regions and renewal outcomes over 6-12 months; causal retention uplift remains unproven.

Ground truth | team only / 7

Capacity and economics

Illustrative scale assumption, not a forecast: 1,000 eligible complex cases/month x 15 net minutes saved / 60 = 250 hours/month. At EUR 45/hour that is EUR 11,250 of capacity value. Reassign 80 hours to proactive account reviews, 70 to reference maintenance and learning, and 100 to complex backlog reduction. Named owners: Customer Success lead, Enablement lead and Support Operations lead respectively. Confirm this allocation in scheduling.

ROI without double counting

Assume EUR 45,000 build plus EUR 3,000/month running cost. Twelve-month total cost is EUR 81,000. Twelve months at full adoption gives EUR 135,000 gross capacity value and an illustrative capacity-valued ROI of $(135,000 - 81,000) / 81,000 = 66.7\%$. Break-even volume is 600 eligible cases/month at the same net minutes. At 500 cases/month the same calculation is -16.7%. Ramp-up, adoption failure and extra maintenance would reduce returns. This is not cash ROI: no labor saving is realized without an actual spend change. Do not add the value of redeployed hours a second time.

Customer value pathway

At a 650-account future opening cohort, a five-percentage-point retention difference represents 32.5 accounts in expectation. At the current average ARR of EUR 180,000/account, that is EUR 5.85 million of illustrative ARR exposure, not attributable revenue or profit. Account mix, contract timing and many non-service factors matter. Do not book this as a pilot benefit or promise the 93% target.

Measure the capacity decision

Ask: If capacity is released, what happens to that capacity? Compare scheduled and actually used hours, proactive conversations completed, backlog age and reused lessons. Hours freed without a changed operating choice may remain local efficiency.

Ground truth | team only / 8

Two-day prototype experiment

Day 1: ingest only the nine evidence files; preserve file names and record IDs. Build the strategy-to-evidence view and a candidate comparison using seven criteria. Reserve AST-0041 to AST-0050 and linked customer/agent comments as a holdout if testing generalization. Day 2: run five evaluation scenarios in fresh sessions, record outputs before showing the key, and review evidence links, alternatives, intervention choices and uncertainty. The source KB remains available so holdout cases can test applicability.

Eight-week operating pilot design

Weeks 1-2: establish a baseline on a permissioned set of 60 complex historical cases and recruit two comparable regional groups. Weeks 3-4: curate reference ownership and scope. Weeks 5-7: counterbalance matched cases between usual practice, governed references alone and governed references plus AI, with no case repeated to the same agent. Stratify product, complexity and tenure; record customer waiting separately. Week 8: review net handling time, answer correctness, escalation appropriateness, agent agency and capacity use. Historical simulation cannot establish live retention impact.

Success and stop criteria

Proposed gates, not established outcomes: at least 95% version-applicability accuracy against specialist review; at least 15% lower active handling time versus usual practice; no worse critical-answer safety; no cross-tenant disclosure or unauthorized action. Check AI adds benefit beyond governed references alone. Stop immediately on disclosure or unauthorized action. Pause if material error exceeds 5%. Keep justified escalations. Report uncertainty and small-sample limits; do not infer churn reduction from this pilot.